


Artificial Intelligence

Introduction



(slides adapted from Anca Dragan)

Robots look like this. We have autonomous cars that figure out how to take us to our destination





Robots help nurses in hospitals deliver stuff to different rooms



Drones that record cool videos of us as we
do outdoor activities

The state of the art of AI

○ [Link](#)





Child 14, happy

Child 13, angry

Child 9, neutral

Child 0, neutral

Child 10, neutral

Child 7, sad

Child 4, neutral

Child 12, neutral

Child 3, neutral

Child 1, neutral

Child 11, neutral




State of the Art: **Human vs Robot**



A reminder for engineers trying to build AI systems that achieve human-level performance. It's often a lot harder than we at first realize. Humans are amazing.



clideo.com



Sofia



Sofia & “indeed” default answer



Deadly orders and rules that Blinly applied

- Rules: The mom said “if you keep doing this crap, I will tell Blinky to cook you for diner”
- Orders: Alex, “kill me, kill my parents, kill everybody

3 Types of Artificial Intelligence

Artificial Narrow Intelligence (ANI)



Stage-1

Machine Learning

- Specialises in one area and solves one problem



Siri



Alexa



Cortana

Artificial General Intelligence (AGI)



Stage-2

Machine Intelligence

- Refers to a computer that is as smart as a human across the board

Artificial Super Intelligence (ASI)



Stage-3

Machine Consciousness

- An intellect that is much smarter than the best human brains in practically every field



Photo: Unsplash

توركا

دولة في الشرق الأوسط

توركيا - دولة الشرق الأوسط
توركيا دولة في الشرق الأوسط
توركيا دولة في الشرق الأوسط
توركيا دولة في الشرق الأوسط



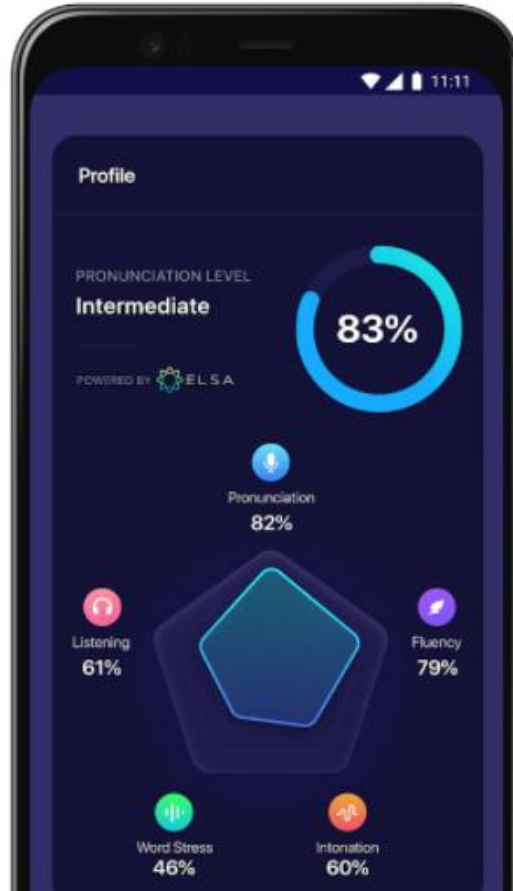
توركيا دولة في الشرق الأوسط

توركيا دولة في الشرق الأوسط

Alexa helps to solve assignments 😊




AI powered Apps (Elsa App)



Today

- What is artificial intelligence?
- Where did it come from/What can AI do?
 - What should we and shouldn't we worry about? What can we do about the things we should worry about?
- What is this course?



What is AI?

The science of making machines that:

Rational Decisions


We'll use the term **rational** in a very specific, technical way:

- Rational: maximally achieving pre-defined goals
- Rationality only concerns what decisions are made
(not the thought process behind them)
- Goals are expressed in terms of the **utility** of outcomes
- Being rational means **maximizing your expected utility**

A better title for this course would be:


Computational Rationality

Maximize Your Expected Utility




What About the Brain?

- Brains (human minds) are very good at making rational decisions, but not perfect
- Brains aren't as modular as software, so hard to reverse engineer!
- “Brains are to intelligence as wings are to flight”
- Lessons learned from the brain: memory and simulation are key to decision making





Designing Rational Agents

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that maximize its (expected) **utility**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions




Pac-Man as an Agent




Textbook

- Russell & Norvig, AI: A Modern Approach




A (Short) History of AI



A (Short) History of AI


- 1940-1950: Early days
 - 1943: McCulloch & Pitts: Boolean circuit model of brain
 - 1950: Turing's "Computing Machinery and Intelligence"
- 1950—70: Excitement: Look, Ma, no hands!
 - 1950s: Early AI programs, including Samuel's checkers program, Newell & Simon's Logic Theorist, Gelernter's Geometry Engine
 - 1956: Dartmouth meeting: "Artificial Intelligence" adopted
 - 1965: Robinson's complete algorithm for logical reasoning
- 1970—90: Knowledge-based approaches
 - 1969—79: Early development of knowledge-based systems
 - 1980—88: Expert systems industry booms
 - 1988—93: Expert systems industry busts: "AI Winter"
- 1990—: Statistical approaches
 - Resurgence of probability, focus on uncertainty
 - General increase in technical depth
 - Agents and learning systems... "AI Spring"?
- 2000—: Where are we now?



What Can AI Do?

Quiz: Which of the following can be done at present?

- ✓ Play a decent game of Jeopardy?
- ✓ Win against any human at chess?
- ✓ Win against the best humans at Go?
- ✓ Play a decent game of tennis?
- ✓ Grab a particular cup and put it on a shelf?
- ✗ Unload any dishwasher in any home?
- ⚡ Drive safely along the highway?
- ✗ Drive safely along Telegraph Avenue?
- ✓ Buy a week's worth of groceries on the web?
- ✗ Buy a week's worth of groceries at Berkeley Bowl?
- ⚡ Discover and prove a new mathematical theorem?
- ✗ Perform a surgical operation?
- ✗ Unload a know dishwasher in collaboration with a person?
- ✓ Translate spoken Chinese into spoken English in real time?
- ✓ Write story?
- ✗ Write an intentionally funny story?





Game Agents

- Classic Moment: May, '97: Deep Blue vs. Kasparov
 - First match won against world champion
 - “Intelligent creative” play
 - 200 million board positions per second
 - Humans understood 99.9 of Deep Blue's moves
 - Can do about the same now with a PC cluster
- 1996: Kasparov Beats Deep Blue

“I could feel --- I could smell --- a new kind of intelligence across the table.”
- 1997: Deep Blue Beats Kasparov

“Deep Blue hasn't proven anything.”




Game Agents

- Reinforcement learning



Pong



Enduro



Robotics

Demo 1: ROBOTICS – soccer.avi



Demo 4: ROBOTICS – laundry.avi

Demo 2: ROBOTICS – soccer2.avi

Demo 5: ROBOTICS – petman.avi

Demo 3: ROBOTICS – gcar.avi


- Robotics
 - Part mech. eng.
 - Part AI
 - Reality much harder than simulations!
- Technologies
 - Vehicles
 - Rescue
 - Help in the home
 - Lots of automation...
- In this class:
 - We ignore mechanical aspects
 - Methods for planning
 - Methods for control



Images from UC Berkeley, Boston Dynamics, RoboCup, Google

Natural Language

- Speech technologies (e.g. Siri)
 - Automatic speech recognition (ASR)
 - Text-to-speech synthesis (TTS)
 - Dialog systems
- Language processing technologies
 - Question answering
 - Machine translation



"Il est impossible aux journalistes de rentrer dans les régions tibétaines"

Bruno Philip, correspondant du "Monde" en Chine, estime que les journalistes de l'AFP qui ont été expulsés de la province tibétaine du Qinghai "n'étaient pas dans l'illégalité".

Les faits Le dalaï-lama dénonce l'"enfer" imposé au Tibet depuis sa fuite, en 1959

Vidéo Anniversaire de la rébellion



"It is impossible for journalists to enter Tibetan areas"

Philip Bruno, correspondent for "World" in China, said that journalists of the AFP who have been deported from the Tibetan province of Qinghai "were not illegal."

Facts The Dalai Lama denounces the "hell" imposed since he fled Tibet in 1959

Video Anniversary of the Tibetan rebellion: China on guard




- Web search
- Text classification, spam filtering, etc...

Computer Vision



"man in black shirt is playing guitar."




"construction worker in orange safety vest is working on road."




"two young girls are playing with lego toy."



"boy is doing backflip on wakeboard."



"girl in pink dress is jumping in air."



"black and white dog jumps over bar."



"young girl in pink shirt is swinging on swing."




"man in blue wetsuit is surfing on wave."

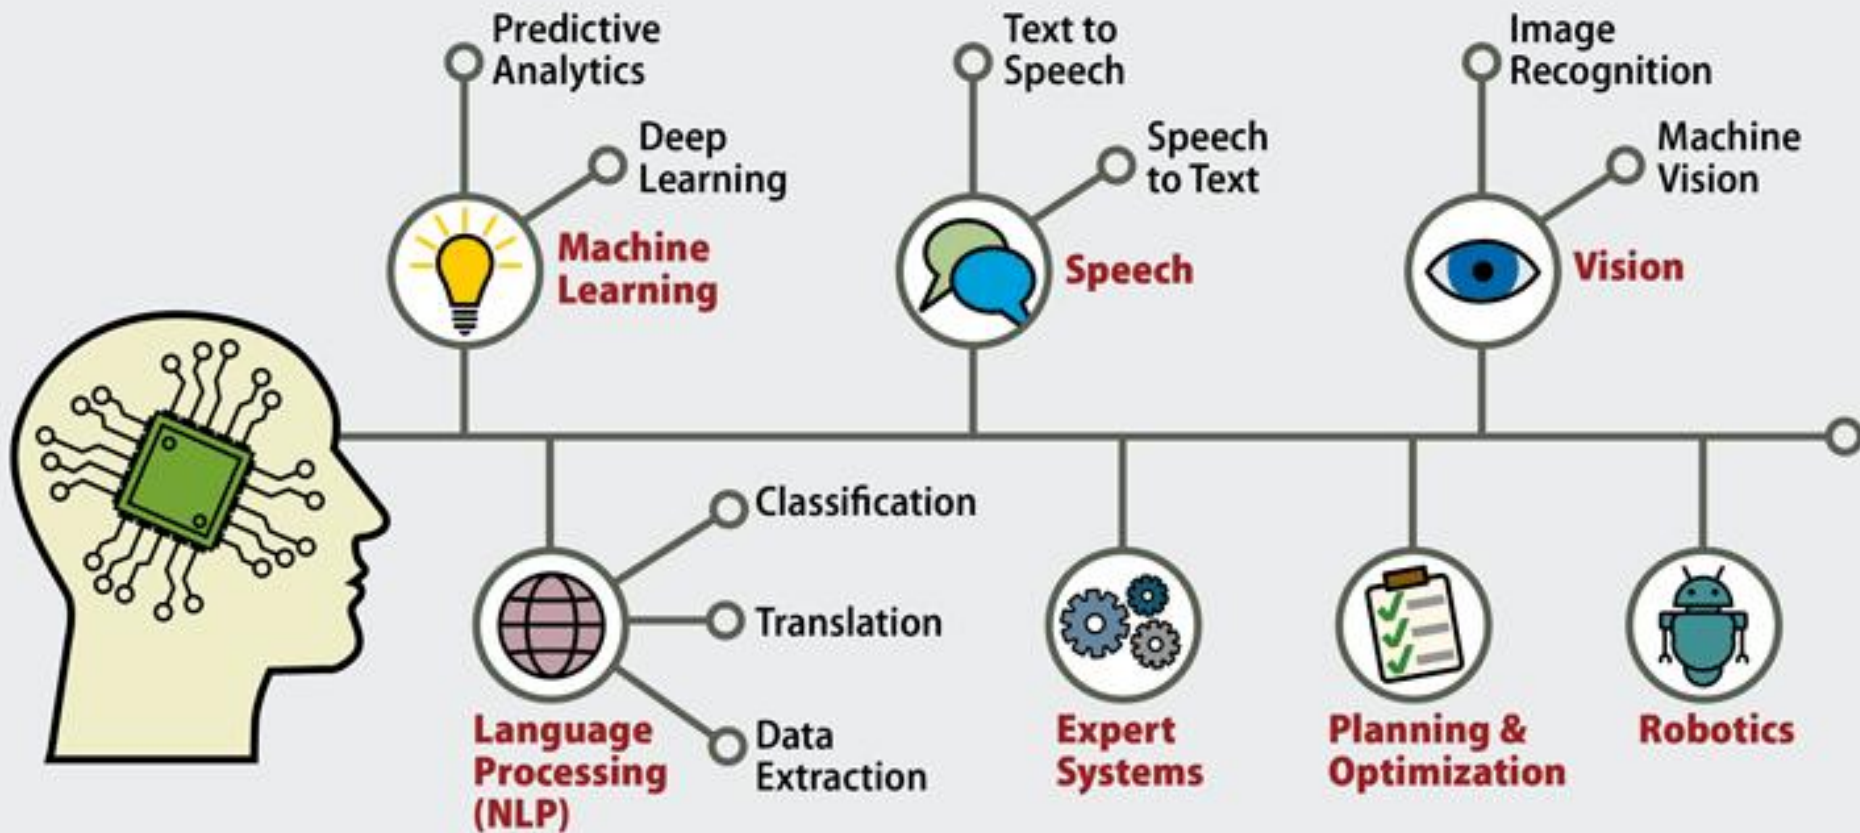
Karpathy & Fei-Fei, 2015; Donahue et al., 2015; Xu et al, 2015; many more

COMPUTER VISION WITH MASK R-CNN





Artificial Intelligence



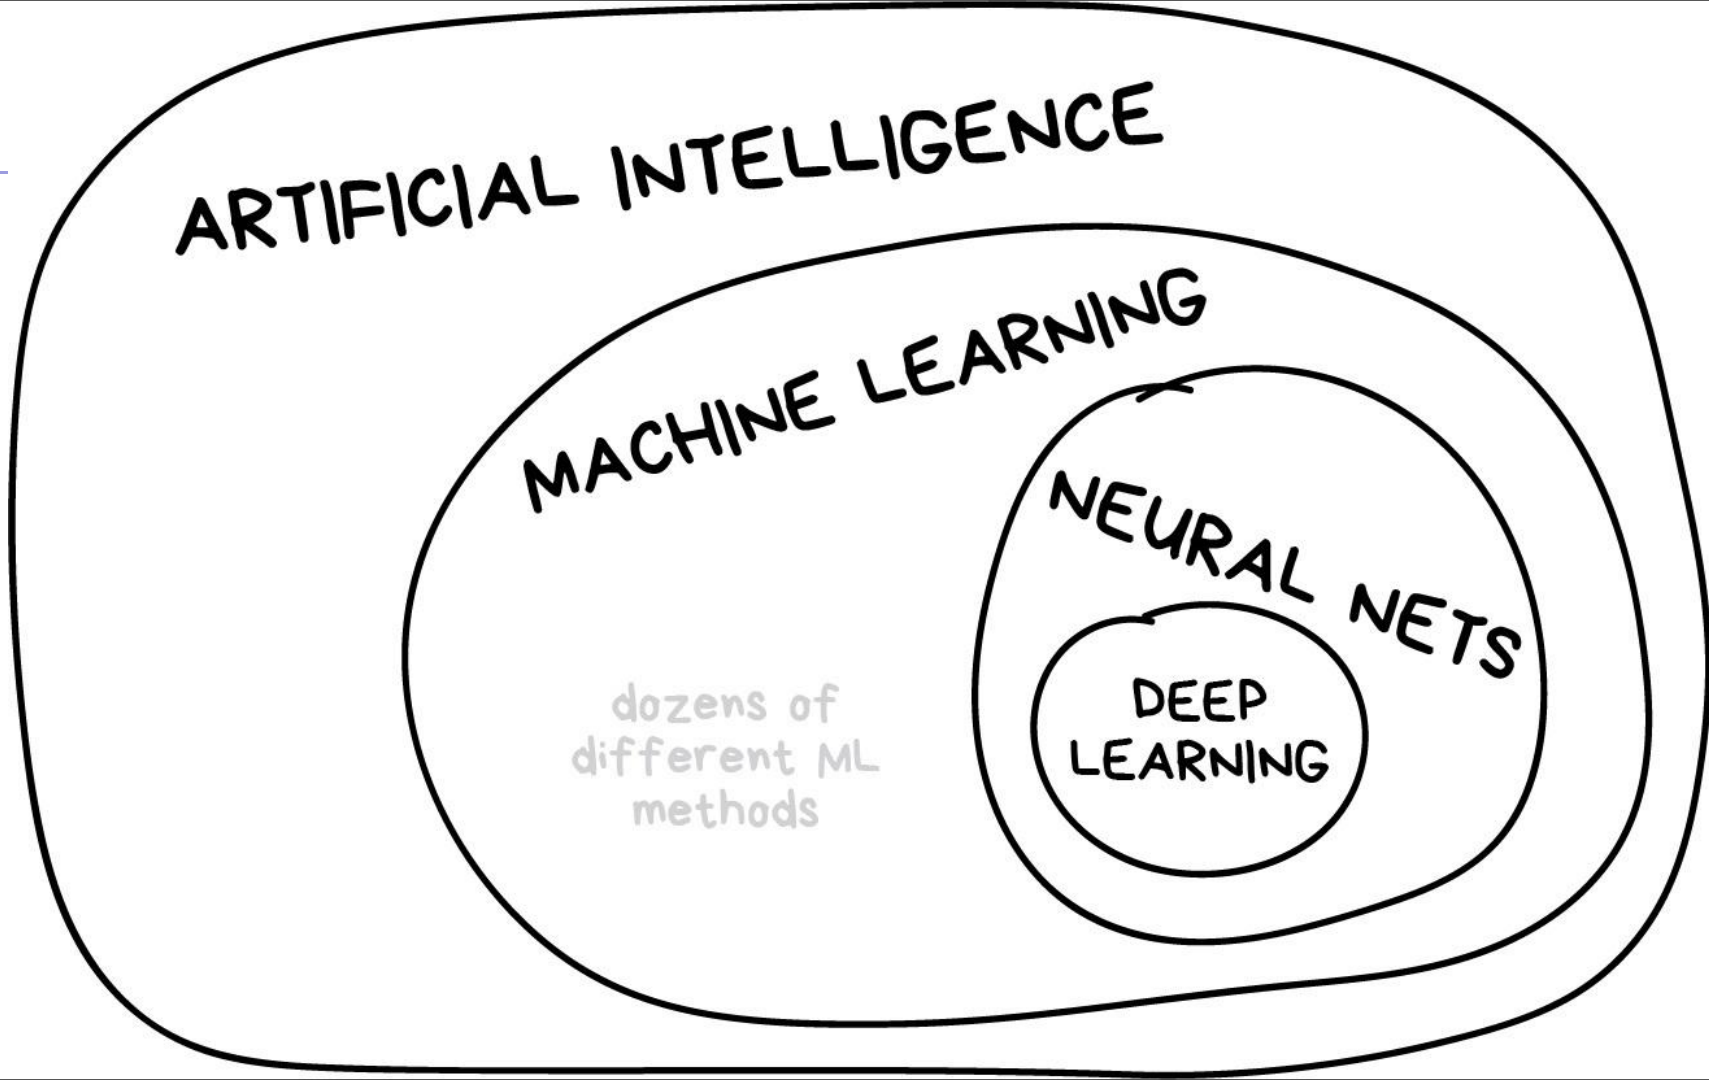
ARTIFICIAL INTELLIGENCE

MACHINE LEARNING

NEURAL NETS

DEEP
LEARNING

dozens of
different ML
methods



CLASSICAL MACHINE LEARNING

Data is pre-categorized
or numerical

SUPERVISED

Predict
a category

CLASSIFICATION

«Divide the socks by color»



Predict
a number

REGRESSION

«Divide the ties by length»



Data is not labeled
in any way

UNSUPERVISED

Divide
by similarity

CLUSTERING

«Split up similar clothing
into stacks»



Identify sequences

Find hidden
dependencies

ASSOCIATION

«Find what clothes I often
wear together»




DIMENSION REDUCTION

(generalization)

«Make the best outfits from the given clothes»





Classification

THE SIMPLEST SPAM-FILTER

(used until 2010)

hey	...	1829
I'm	...	1710
no	...	1191
where	...	1012
you	...	985
speak	...	873
learn	...	747
one	...	739

good letters

viagra	...	1552
casino	...	1492
100%	...	1320
credit	...	1184
sale	...	985
press	...	873
free	...	747
enlarge	...	739

spam letters

672 times

«KITTY»


13 times

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

BAYES' THEOREM

NOT SPAM

NAIVE BAYES



Clustering

